



计算机视觉

邬向前

计算学部

多模态智能及应用研究中心

电子邮箱: xqw@hit.edu.cn

Lecture 2:

Image Filtering

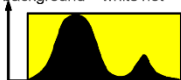
Outline

- **Point (Pixel) Operations**
- **Group (Neighborhood) Operations**



Image Histograms

Light object(s) against a darker background - 'white hot'



Dark object(s) against a lighter background - 'black hot'



Bright
object

Dark
background

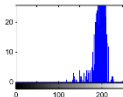
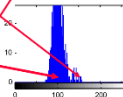
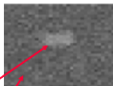


Image Histograms

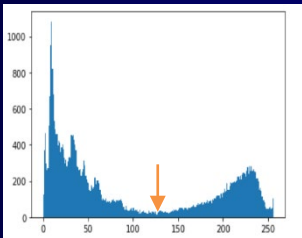
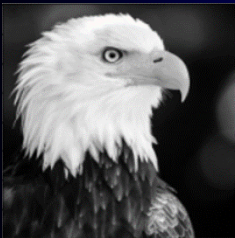
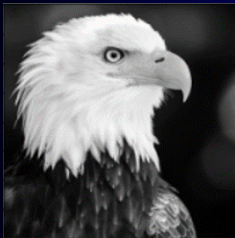


Image Histograms



上堂课的作业：

1. 打开任意一幅灰度图片，计算该图像的图像直方图；
2. 打开任意一幅灰度图片，对该图像进行处理，使得整幅图像变亮10个灰度级；
3. 打开任意一幅灰度图片，在该图像的中心位置画一个半径为半个图像高度或宽度的圆。

Point (Pixel) Operations

Examples of point processing

How would you implement these?

original



darken



lower contrast



non-linear lower contrast



\mathcal{I}

invert



lighten



raise contrast



non-linear raise contrast



Examples of point processing

How would you implement these?

original



$$x$$

darken



$$x - 128$$

lower contrast



$$\frac{x}{2}$$

non-linear lower contrast



$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

lighten



$$x + 128$$

raise contrast



$$x \times 2$$

non-linear raise contrast



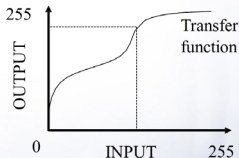
$$\left(\frac{x}{255}\right)^2 \times 255$$

Point (Pixel) Operation

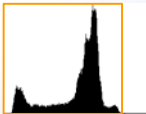
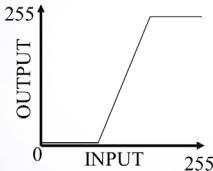
- Point operation: A function is applied to every pixel in an image, which operates only on the pixel's current value.
- Thresholding - A mask may be created by setting a pixel value to 1 or 0 depending upon if the current value is above or below a certain threshold value.

Input pixel value, I , mapped to output pixel value, O , via transfer function T .

$$O = T(I)$$



Contrast Enhancement: Linear Stretching

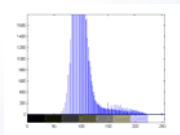
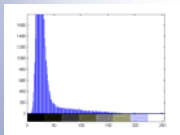
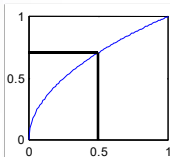


Contrast Enhancement: Power Law Function

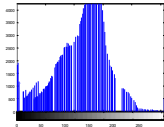
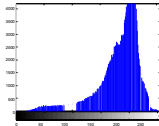
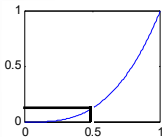
$$O = I^\gamma$$

- $\gamma < 1$ to enhance contrast in dark regions
- $\gamma > 1$ to enhance contrast in bright regions.

Contrast Enhancement: Power Law Function ($\gamma = 0.5$)



Contrast Enhancement: Power Law Function ($\gamma = 3.0$)



Contrast Enhancement: Power Law Function

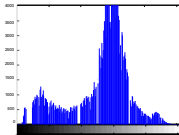
$$\gamma=2$$

- Look-up Table
 - Transfer function implemented as a look-up table (LUT).
 - Implemented in hardware or software.

<i>I</i>	<i>O</i>
0	0
0.10	0.01
0.20	0.04
0.30	0.09
0.40	0.16
0.50	0.25
0.60	0.36
0.70	0.49
0.80	0.64
0.90	0.81
1.00	1.00

Contrast Enhancement: Histogram Equalisation

- Image histograms consisting of peaks and low plains.
- Peaks = many pixels concentrated in a few grey levels
- Plains = small number of pixels distributed over a wider range of grey levels



Contrast Enhancement: Histogram Equalisation

■ 直方图均衡化的理论基础

- 前提：如果一幅图像占有全部可能的灰度级，并且均匀分布
- 结论：该图像具有高对比度和多变的灰色色调
- 外观：图像细节丰富，质量更高

Contrast Enhancement: Histogram Equalisation

■ 直方图均衡化的步骤

1. 计算累计直方图
2. 将累计直方图进行区间转换
3. 在累计直方图中，概率相近的原始值，会被处理为相同的值。

Contrast Enhancement: Histogram Equalisation

■ 直方图均衡化的步骤

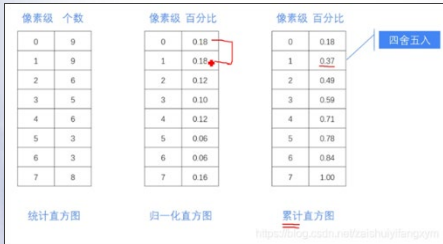
1. 计算累计直方图

原始图像							像素级 个数	
0	1	4	1	7	3	3	0	9
0	0	4	0	0	1	3	1	9
1	2	7	5	7	4	6	2	6
0	4	0	1	1	6	6	3	5
7	1	2	2	7	3	3	4	6
4	5	7	4	2	7	2	5	3
0	7	1	5	2	0	1	6	3
							7	8
统计直方图								

Contrast Enhancement: Histogram Equalisation

■ 直方图均衡化的步骤

1. 计算累计直方图



Contrast Enhancement: Histogram Equalisation

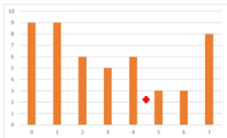
- 直方图均衡化的步骤
- 2. 将累计直方图进行区间转换
- 3. 在累计直方图中，概率相近的原始值，会被处理为相同的值。



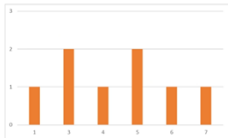
Contrast Enhancement: Histogram Equalisation

■ 直方图均衡化的步骤

2. 将累计直方图进行区间转换

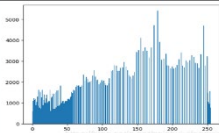
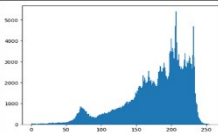


原始直方图



均衡直方图

Histogram Equalisation



Contrast Enhancement: Comparison

Original



$$\gamma > 1$$



Histogram
equalisation

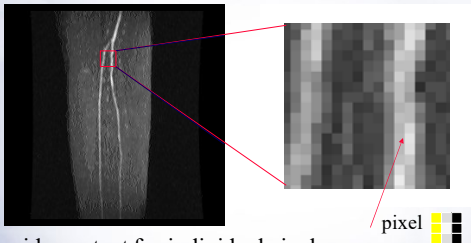


Group (Neighborhood) Operations

Neighbourhood Operations

- Replace each pixel by a *linear* combination of its neighbors (and possibly itself).
- The combination is determined by the filter's *kernel*.
- The same kernel is *shifted* to all pixel locations so that all pixels use the same linear combination of their neighbors.

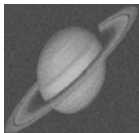
Why are Neighbourhoods Important?



- Provide context for individual pixels.
- Relationships between neighbours determine image features.

Neighbourhood Operations

Noise Reduction



Edge Enhancement



Zooming



Convolution

Convolution for 1D continuous signals

Definition of filtering as convolution:

filtered signal $\Rightarrow (f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x-y)dy$

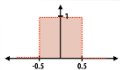
notice the flip

filter input signal

Consider the box filter example:

1D continuous
box filter

$$f(x) = \begin{cases} 1 & |x| \leq 0.5 \\ 0 & \text{otherwise} \end{cases}$$



filtering output is a
blurred version of
g

$$(f * g)(x) = \int_{-0.5}^{0.5} g(x-y)dy$$

连续卷积的通俗解释：做馒头

楼下早点铺子生意太好了，供不应求，就买了一台机器，不断的生产馒头。

假设馒头的生产速度是 $f(t)$ ，那么一天后生产出来的馒头总量为：

$$\int_0^{24} f(t) dt$$

馒头生产出来之后，就会慢慢腐败，假设腐败函数为 $g(t)$ ，比如，10个馒头，24小时会腐败：

$$10 * g(t)$$

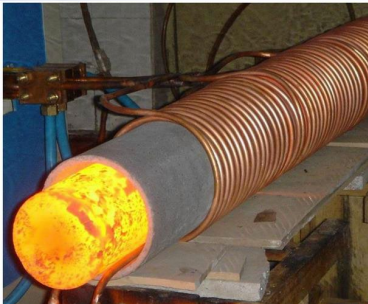
想想就知道，第一个小时生产出来的馒头，一天后会经历24小时的腐败，第二个小时生产出来的馒头，一天后会经历23小时的腐败。

如此，我们可以知道，一天后，馒头总共腐败了：

$$\int_0^{24} f(t)g(24 - t)dt$$

这就是连续的卷积。

连续卷积的通俗解释：铁棒加热



$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i, j = -\infty}^{\infty} f(i, j) I(x - i, y - j)$$

Diagram annotations for the equation above:

- An arrow points from the text "filtered image" to $(f * g)(x, y)$.
- An arrow points from the text "notice the flip" to $I(x - i, y - j)$.
- An arrow points from the text "filter" to $f(i, j)$.
- An arrow points from the text "input image" to $I(x - i, y - j)$.

If the filter $f(i, j)$ is non-zero only within $-1 \leq i, j \leq 1$, then

$$(f * g)(x, y) = \sum_{i, j = -1}^1 f(i, j) I(x - i, y - j)$$

The kernel we saw earlier is the 3x3 matrix representation of $f(i, j)$.

In general, the kernel for image processing is symmetry, so the flip processing can be skipped.

Convolution

- Consists of filtering an image A using a filter (mask) B .
- Mask is a small image whose pixel values are called weights.
- Weights modify relationships between pixels.

Filter,
mask or
template **B**

$B_{1,1}$	$B_{1,2}$
$B_{2,1}$	$B_{2,2}$

2×2

Input
image **A**

$A_{1,1}$	$A_{1,2}$	$A_{1,3}$	$A_{1,4}$
$A_{2,1}$	$A_{2,2}$	$A_{2,3}$	$A_{2,4}$
$A_{3,1}$	$A_{3,2}$	$A_{3,3}$	$A_{3,4}$
$A_{4,1}$	$A_{4,2}$	$A_{4,3}$	$A_{4,4}$

4×4

Convolved
Image **C**

$C_{1,1}$	$C_{1,2}$	$C_{1,3}$
$C_{2,1}$	$C_{2,2}$	$C_{2,3}$
$C_{3,1}$	$C_{3,2}$	$C_{3,3}$

3×3

Convolution

$BA_{1,1}$	$BA_{1,2}$
$BA_{2,1}$	$BA_{2,2}$

$A_{1,3}$	$A_{1,4}$
$A_{2,3}$	$A_{2,4}$

$A_{3,1}$	$A_{3,2}$	$A_{3,3}$	$A_{3,4}$
$A_{4,1}$	$A_{4,2}$	$A_{4,3}$	$A_{4,4}$

$A_{1,1} \times B_{1,1}$	$A_{1,2} \times B_{1,2}$
$A_{2,1} \times B_{2,1}$	$A_{2,2} \times B_{2,2}$

$$C_{1,1} = \boxed{A_{1,1} \times B_{1,1}} + \boxed{A_{1,2} \times B_{1,2}} + \boxed{A_{2,1} \times B_{2,1}} + \boxed{A_{2,2} \times B_{2,2}}$$

Convolution

$A_{1,1}$	$A_{1,2}$	$BA_{1,3}$	$BA_{1,4}$
$A_{2,1}$	$A_{2,2}$	$BA_{2,3}$	$BA_{2,4}$
$A_{3,1}$	$A_{3,2}$	$A_{3,3}$	$A_{3,4}$
$A_{4,1}$	$A_{4,2}$	$A_{4,3}$	$A_{4,4}$

$$A_{1,3} \times B_{1,1} \quad A_{1,4} \times B_{1,2}$$

$$A_{2,3} \times B_{2,1} \quad A_{2,4} \times B_{2,2}$$

$$C_{1,3} = \boxed{A_{1,3} \times B_{1,1}} + \boxed{A_{1,4} \times B_{1,2}} + \boxed{A_{2,3} \times B_{2,1}} + \boxed{A_{2,4} \times B_{2,2}}$$

Convolution

$A_{1,1}$	$A_{1,2}$	$A_{1,3}$	$A_{1,4}$
$BA_{2,1}$	$BA_{2,2}$	$A_{2,3}$	$A_{2,4}$
$BA_{3,1}$	$BA_{3,2}$	$A_{3,3}$	$A_{3,4}$
$A_{4,1}$	$A_{4,2}$	$A_{4,3}$	$A_{4,4}$

$$A_{2,1} \times B_{1,1} \quad A_{2,2} \times B_{1,2}$$

$$A_{3,1} \times B_{2,1} \quad A_{3,2} \times B_{2,2}$$

$$C_{2,1} = \boxed{A_{2,1} \times B_{1,1}} + \boxed{A_{2,2} \times B_{1,2}} + \boxed{A_{3,1} \times B_{2,1}} + \boxed{A_{3,2} \times B_{2,2}}$$

Mathematical Notation

$$C_{1,1} = \boxed{A_{1,1} \times B_{1,1}} + \boxed{A_{1,2} \times B_{1,2}} + \boxed{A_{2,1} \times B_{2,1}} + \boxed{A_{2,2} \times B_{2,2}}$$

$$B = M \times N$$

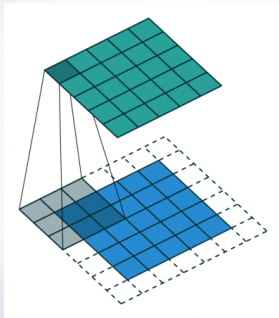
$$C_{k,l} = \sum_{i=k}^{k+M-1} \sum_{j=l}^{l+N-1} A_{i,j} \times B_{i-k+1,j-l+1}$$

$$\sum_{i=1}^2 i = 1 + 2$$

$$\sum_{i=1}^2 A_i = A_1 + A_2$$

$$\sum_{i=1}^2 \sum_{j=1}^2 A_{i,j} = \sum_{i=1}^2 (A_{i,1} + A_{i,2}) = A_{1,1} + A_{1,2} + A_{2,1} + A_{2,2}$$

Convolution



Convolution

Filter,
mask or
template

B

-1	2
-1	2

2×2

Input
image

A

4	4	7	9
4	3	8	9
3	5	9	9
3	6	10	9

4×4

Convolved
Image

C

6	23	21
9	26	19
16	27	17

3×3

=

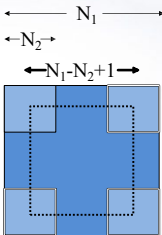
Convolution Size

Image size = $M_1 \times N_1$

Mask size = $M_2 \times N_2$

Convolution size =
 $(M_1 - M_2 + 1) \times (N_1 - N_2 + 1)$

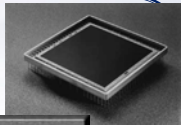
Typical Mask sizes
 $= 3 \times 3, 5 \times 5, 7 \times 7,$
 $9 \times 9, 11 \times 11$



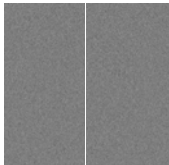
What is the convolved image size for a 128×128 image and 7×7 mask?

Noise

- Source of noise = CCD chip.
- Electronic signal fluctuations in detector.
- Caused by thermal energy.
- Worse for infra-red sensors.



image



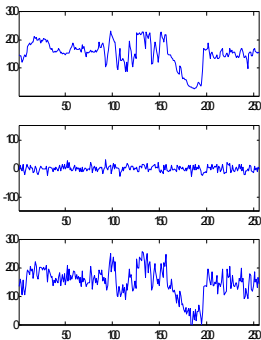
noise



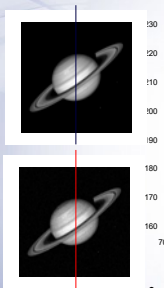
'grainy' image

Noise

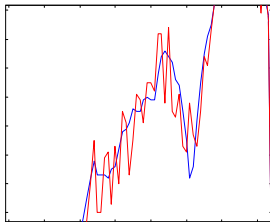
- Plot of image brightness.
- Vertical slice through image.
- Noise is additive.
- Noise fluctuations are rapid, i.e, high frequency.



Noise Reduction



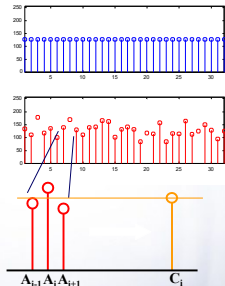
130
120
110
100
90
80
180
170
160



- Noise varies above and below uncorrupted image.

Noise Reduction-1st principles

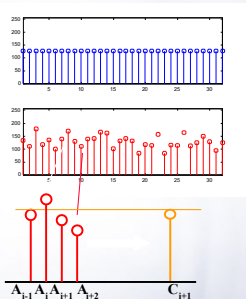
- How do we reduce noise?
- Consider a uniform 1-d image A and add noise.
- Focus on a pixel neighbourhood.
- Central pixel has been increased and neighbouring pixels have decreased.



Noise Reduction-1st principles

- Averaging ‘smoothes’ the noise fluctuations.
- Consider the next pixel A_{i+1}
- Repeat for remainder of pixels.

$$C_{i+1} = \frac{A_i + A_{i+1} + A_{i+2}}{3}$$



Noise Reduction- Neighborhood Operations

- All pixels can be averaged by convolving 1-d image A with mask B to give enhanced image C .
- Weights of B must equal one when added together.

$$C = A * B$$

$$B = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$C = A * B$$

$$B = [B_1 \quad B_2 \quad B_3]$$

$$C_i = A_{i-1} \times B_1 + A_i \times B_2 + A_{i+1} \times B_3$$

$$B = \frac{1}{3} [1 \quad 1 \quad 1]$$

$$C_i = \frac{A_{i-1} + A_i + A_{i+1}}{3}$$

- Extend to two dimensions.

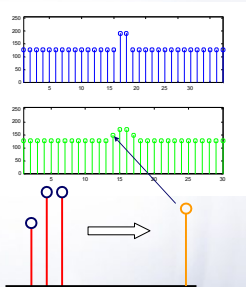
Noise Reduction

- Technique relies on high frequency noise fluctuations being 'blocked' by filter. Hence, low-pass filter.
- Fine detail in image may also be smoothed.
- Balance between keeping image fine detail and reducing noise.
- Example:
 - Saturn image coarse detail
 - Boat image contains fine detail
 - Noise reduced but fine detail also smoothed



Noise Reduction

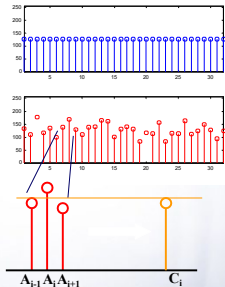
- Consider a uniform 1-d image A with a step function.
- Step function corresponds to fine image detail such as an edge.
- Low-pass filter ‘blurs’ the edge.



Noise Reduction-1st principles

- How do we reduce noise without averaging?
- Consider a uniform 1-d image A and add noise.
- Focus on a pixel neighbourhood.
- Non-linear operator?

Median filter!



Noise Reduction- Neighborhood Operations

- All pixels can be replaced by neighbourhood median by convolving 1-d image A with median filter B to give enhanced image C.

$$\mathbf{C} = \mathbf{A} * \mathbf{B}$$

$$\mathbf{B} = [B_1 \quad B_2 \quad B_3]$$

$$C_i = \text{median}\{A_{i-1} \times B_1, A_i \times B_2, A_{i+1} \times B_3\}$$

$$\mathbf{B} = [1 \quad 1 \quad 1]$$

$$C_i = \text{median}\{A_{i-1}, A_i, A_{i+1}\}$$
- Extend to two dimensions.

$$C_{k,l} = \text{median}_{i=k:M-1, j=l:N-1} \{A_{i,j} \times B_{i-k+1, j-l+1}\}$$

$$B_{i,j} = 1 \text{ for all } i,j$$

Noise reduction

Original



Low-pass

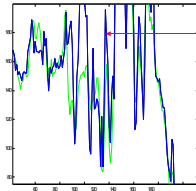


Median

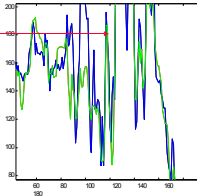


Noise reduction

Low-pass



Median



- Low-pass: fine detail smoothed by averaging
- Median: fine detail passed by filter

Filter Operations

1 Low-Pass Filter

- Class: Image Enhancement/Restoration
- Implementation: Pixel group process and smooth an image

1 1 1

1 1 1

1 2 1

1 1 1

1 2 1

2 4 2

1 1 1

1 1 1

1 2 1

/9

/10

/16

2 High-Pass Filter

- Implementation: Pixel group process and sharpen an image

-1 -1 -1

0 -1 0

1 -2 1

-1 9 -1

-1 5 -1

-2 5 -2

-1 -1 -1

0 -1 0

1 -2 1

Filter Operations

3 The Gaussian filter

- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$



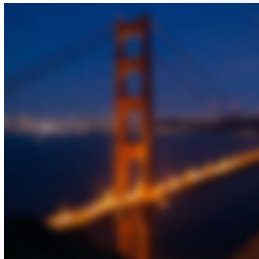
Is this a separable filter?

- weight falls off with distance from center pixel
- theoretically infinite, in practice truncated to some maximum distance
- Any heuristics for selecting where to truncate?
- usually at $2-3\sigma$

kernel $\frac{1}{16}$

1	2	1
2	4	2
1	2	1

Gaussian filtering example



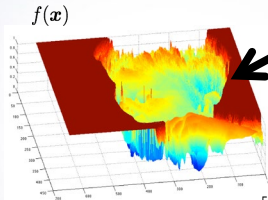
Edge Enhancement Neighborhood Operations

- Sobel Filter
- Derivative of Gaussian (DoG) filter
- Derivative of Laplacian (LoG) filter

What are image edges?



grayscale image



Very sharp
discontinuities
in intensity.

domain $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

Detecting edges

How would you go about detecting edges in an image (i.e., discontinuities in a function)?

✓ You take derivatives: derivatives are large at discontinuities.

How do you differentiate a discrete image (or any other discrete signal)?

✓ You use finite differences.

Finite differences

High-school reminder: definition of a derivative using forward difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

Alternative: use central difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+0.5h) - f(x-0.5h)}{h}$$

For discrete signals: Remove limit and set $h = 2$

$$f'(x) = \frac{f(x+1) - f(x-1)}{2}$$

1D derivative filter

-1	0	1
----	---	---

1	0	-1
---	---	----

The Sobel filter

Horizontal Sober filter:

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

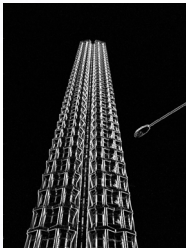
Vertical Sober filter:

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Sobel filter example



original

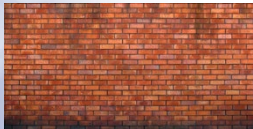


horizontal Sobel filter

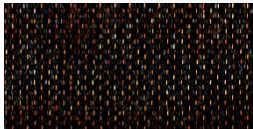


vertical Sobel filter

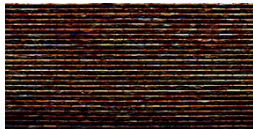
Sobel filter example



original



horizontal Sobel filter



vertical Sobel filter

Computing image gradients

1. Select your favorite derivative filters.

$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$S_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

2. Convolve with the image to compute derivatives.

$$\frac{\partial f}{\partial x} = S_x \otimes f$$

$$\frac{\partial f}{\partial y} = S_y \otimes f$$

3. Form the image gradient, and compute its direction and amplitude.

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

gradient

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

direction

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

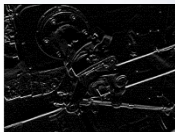
amplitude

Image gradient example

original



vertical
derivative



gradient
amplitude



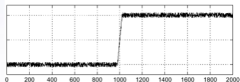
horizontal
derivative



How does the gradient direction relate to these edges?

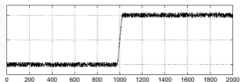
How do you find the edge of this signal?

intensity plot



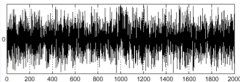
How do you find the edge of this signal?

intensity plot



Using a derivative filter:

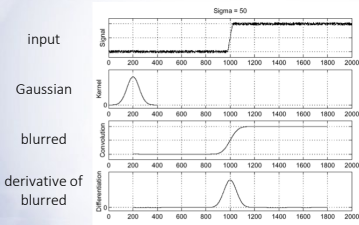
derivative plot



What's the problem here?

Differentiation is very sensitive to noise

When using derivative filters, it is critical to blur first!

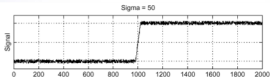


How much
should we blur?

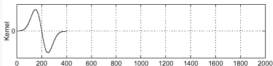
Derivative of Gaussian (DoG) filter

Derivative theorem of convolution: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$

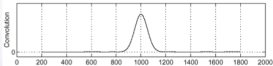
input



derivative of
Gaussian



output (same
as before)



- How many operations did we save?
- Any other advantages beyond efficiency?

Laplace filter

Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order
finite difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$



1D derivative filter

1	0	-1
---	---	----

second-order
finite difference

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$$

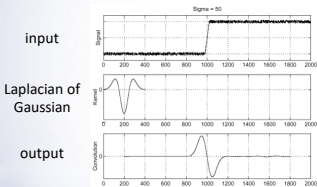


Laplace filter

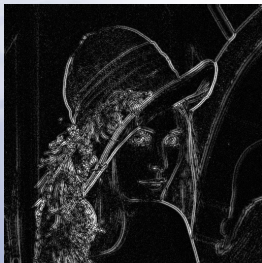
1	-2	1
---	----	---

Laplacian of Gaussian (LoG) filter

As with derivative, we can combine Laplace filtering with Gaussian filtering



Laplace and LoG filtering examples

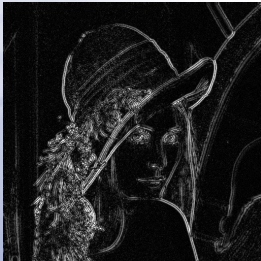


Laplacian of Gaussian filtering



Laplace filtering

Laplacian of Gaussian vs Derivative of Gaussian

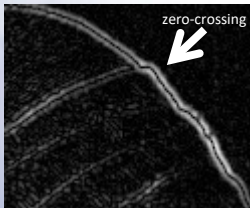


Laplacian of Gaussian filtering

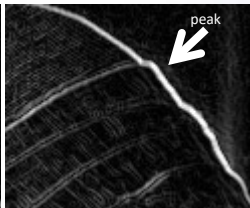


Derivative of Gaussian filtering

Laplacian of Gaussian vs Derivative of Gaussian



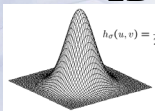
Laplacian of Gaussian filtering



Derivative of Gaussian filtering

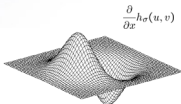
Zero crossings are more accurate at localizing edges (but not very convenient).

2D Gaussian based filters



$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

Gaussian



$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Derivative of Gaussian (DOG)



$$\nabla^2 h_{\sigma}(u, v)$$

Laplacian of Gaussian
(LOG)

作业:

1. 编程实现图像的卷积操作函数;
2. 给定一方差, 编制三个函数, 分别生成高斯滤波器、DOG和LOG滤波器
3. 打开任意一幅灰度图片, 用上面的卷积函数和生成的滤波器, 对其进行滤波并显示

So much for today!



Thank you !!!