



计算机视觉

邬向前

计算学部

多模态智能及应用研究中心

电子邮箱: xqw@hit.edu.cn


Lecture 7 – Optical Flow

Overview

- Segmentation in Video
- Optical flow


Video

- A video is a sequence of frames captured over time
- Now our image data is a function of space (x, y) and time (t)



Applications of segmentation to video

- Background subtraction
 - A static camera is observing a scene
 - Goal: separate the static *background* from the moving *foreground*



Applications of segmentation to video

- Background subtraction
 - Form an initial background estimate

Calculate background: Average a series of preceding images.

$$\frac{1}{N} \sum_{i=1}^N V(x, y, t - i)$$

Background image at time t:

$$B(x, y, t) = \frac{1}{N} \sum_{i=1}^N V(x, y, t - i)$$

Applications of segmentation to video

- Background subtraction

N imatges



Mediana per cada píxel



Model de fons de referència



Applications of segmentation to video

- Background subtraction
 - Form an initial background estimate

$$B(x, y, t) = \frac{1}{N} \sum_{i=1}^N V(x, y, t - i)$$

- Subtract the background estimate from the frame

$$V(x, y, t) - B(x, y, t)$$

- Label as foreground **each pixel** where the magnitude of the difference $> Th$


$$|V(x, y, t) - B(x, y, t)| > Th$$

Applications of segmentation to video

- Background subtraction
 - Form an initial background estimate
 - For each frame:
 - Update estimate using a moving average
 - Subtract the background estimate from the frame
 - Label as foreground each pixel where the magnitude of the difference is greater than some threshold
 - Use median filtering to “clean up” the results

Applications of segmentation to video

- Background subtraction




Applications of segmentation to video

- Background subtraction
 - Pros:
 - Simple;
 - To some extent, overcomes the influence of environmental light;
 - Cons:
 - Can not be used for moving Cameras;
 - Difficult to update the background image in real time


Applications of segmentation to video

- Background subtraction
- Shot boundary detection
 - Commercial video is usually composed of *shots* or sequences showing the same objects or scene
 - Goal: segment video into shots for summarization and browsing (each shot can be represented by a single keyframe in a user interface)
 - Difference from background subtraction: the camera is not necessarily stationary




Applications of segmentation to video

- Background subtraction
- Shot boundary detection



a sudden transition from one shot to another

Applications of segmentation to video




$$Z_n = \Psi(f_n)$$

$$d(f_n, f_{n+1}) = \left(\sum_{k=1} |Z_n(k) - Z_{n+1}(k)|^p \right)^{1/p}$$

f_n : video frame, Ψ : feature extraction function

Z_n : extracted features of video frame f_n



Threshold


Applications of segmentation to video

- Background subtraction
- Shot boundary detection
 - For each frame
 - Compute the distance between the current frame and the previous one
 - » Mean absolute differences (MAD)

Similarity between two consecutive frames pixels

from (i, j) search

$m * n$



$S(x, y)$



$M * N$

$T(x, y)$

Measure of similarity:

$$D(i, j) = \frac{1}{M \times N} \sum_{s=1}^M \sum_{t=1}^N |S(i + s - 1, j + t - 1) - T(s, t)|$$

$$1 \leq i \leq m - M + 1 \quad 1 \leq j \leq n - N + 1$$

Applications of segmentation to video

- Background subtraction
- Shot boundary detection
 - For each frame
 - Compute the distance between the current frame and the previous one
 - » Histogram differences (HD)

Similarity between the histograms of two consecutive frames

$$d(H_1, H_2) = \frac{\sum_I (H_1(I) - \bar{H}_1)(H_2(I) - \bar{H}_2)}{\sqrt{\sum_I (H_1(I) - \bar{H}_1)^2 \sum_I (H_2(I) - \bar{H}_2)^2}}$$

https://blog.csdn.net/qj_33287871

$$\bar{H}_k = \frac{1}{N} \sum_J H_k(J)$$

Correlation: $r(X, Y) = \frac{Cov(X, Y)}{\sqrt{Var[X] Var[Y]}}$

[s://blog.csdn.net/yixinxin](https://blog.csdn.net/yixinxin)

$d(H_1, H_2) \rightarrow 1$, strong Correlation.

Applications of segmentation to video

- Background subtraction
- Shot boundary detection
 - For each frame
 - Compute the distance between the current frame and the previous one
 - » Mean absolute differences (MAD)
 - » Histogram differences (HD)
 - If the distance is greater than some threshold, classify the frame as a shot boundary



Applications of segmentation to video

- Background subtraction
- Shot boundary detection
- Motion Segmentation

Overview

- Segmentation in Video
- Optical flow
 - Lucas-kanade
 - Horn-schunck
 - Optical flow pyramid

Motion estimation: Optical flow



Will start by estimating motion of each pixel separately
Then will consider motion of entire image


Why estimate motion?

Lots of uses


- Track object behavior
- Correct for camera jitter (stabilization)
- Align images (mosaics)
- 3D shape reconstruction
- Special effects



Problem definition: optical flow



$H(x, y)$



$I(x, y)$

How to estimate pixel motion from image H to image I ?


- Solve pixel correspondence problem
 - given a pixel in H , look for **nearby** pixels of the **same color** in I

Key assumptions

- **color constancy**: a point in H looks the same in I
 - For grayscale images, this is brightness constancy
- **small motion**: points do not move very far

This is called the optical flow problem

Optical flow constraints (grayscale images)



Let's look at these constraints more closely

- brightness constancy: Q: what's the equation?

$$H(x, y) = I(x + u, y + v)$$

- small motion: (u and v are less than 1 pixel)
 - suppose we take the Taylor series expansion of I :

$$\begin{aligned} I(x + u, y + v) &= I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \text{higher order terms} \\ &\approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v \end{aligned}$$

Optical flow equation

Combining these two equations

$$\begin{aligned}0 &= I(x + u, y + v) - H(x, y) & I_x &= \frac{\partial I}{\partial x} \\&\approx I(x, y) + I_x u + I_y v - H(x, y) \\&\approx (I(x, y) - H(x, y)) + I_x u + I_y v \\&\approx I_t + I_x u + I_y v \\&\approx I_t + \nabla I \cdot [u \ v]\end{aligned}$$

In the limit as u and v go to zero, this becomes exact

$$0 = I_t + \nabla I \cdot \left[\frac{\partial x}{\partial t} \ \frac{\partial y}{\partial t} \right]$$

Lucas-kanade

Optical flow equation

$$0 = I_t + \nabla I \cdot [u \ v]$$

Q: how many unknowns and equations per pixel?

2 unknowns, one equation

Intuitively, what does this constraint mean?

- The component of the flow in the gradient direction is determined
- The component of the flow parallel to an edge is unknown

This explains the Barber Pole illusion

http://www.sandlotscience.com/Ambiguous/Barberpole_Illusion.htm

<http://www.liv.ac.uk/~marcob/Trieste/barberpole.html>



Solving the aperture problem

How to get more equations for a pixel?

- Basic idea: impose **additional constraints**
 - most common is to assume that the flow field is **smooth locally**
 - one method: **pretend the pixel's neighbors have the same (u,v)**
 - » If we use a 5x5 window, that gives us 25 equations per pixel!

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{array}{ccc} \left[\begin{array}{cc} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{array} \right] & \left[\begin{array}{c} u \\ v \end{array} \right] & = - \left[\begin{array}{c} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{array} \right] \\ \underset{25 \times 2}{A} & \underset{2 \times 1}{d} & \underset{25 \times 1}{b} \end{array}$$

RGB version

How to get more equations for a pixel?

- Basic idea: impose additional constraints
 - most common is to assume that the flow field is smooth locally
 - one method: pretend the pixel's neighbors have the same (u,v)
 - » If we use a 5x5 window, that gives us 25*3 equations per pixel!

$$0 = I_t(\mathbf{p}_i)[0, 1, 2] + \nabla I(\mathbf{p}_i)[0, 1, 2] \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1)[0] & I_y(\mathbf{p}_1)[0] \\ I_x(\mathbf{p}_1)[1] & I_y(\mathbf{p}_1)[1] \\ I_x(\mathbf{p}_1)[2] & I_y(\mathbf{p}_1)[2] \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25})[0] & I_y(\mathbf{p}_{25})[0] \\ I_x(\mathbf{p}_{25})[1] & I_y(\mathbf{p}_{25})[1] \\ I_x(\mathbf{p}_{25})[2] & I_y(\mathbf{p}_{25})[2] \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1)[0] \\ I_t(\mathbf{p}_1)[1] \\ I_t(\mathbf{p}_1)[2] \\ \vdots \\ I_t(\mathbf{p}_{25})[0] \\ I_t(\mathbf{p}_{25})[1] \\ I_t(\mathbf{p}_{25})[2] \end{bmatrix}$$

A
 75×2

d
 2×1

b
 75×1

Note that RGB is not enough to disambiguate because R, G & B are correlated

Lukas-Kanade flow

Prob: we have more equations than unknowns

$$\begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 \quad 25 \times 1 \end{matrix} \longrightarrow \text{minimize } \|Ad - b\|^2$$

Solution: solve least squares problem


- minimum least squares solution given by solution (in d) of:

$$\begin{matrix} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 \quad 2 \times 1 \end{matrix}$$

$$\begin{matrix} \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} & \begin{bmatrix} u \\ v \end{bmatrix} & = & - & \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix} \\ A^T A & \downarrow & & & A^T b \\ & \text{To solve } d=[u,v] & & & \end{matrix}$$

- The summations are over all pixels in the K x K window
- This technique was first proposed by Lukas & Kanade (1981)

Aperture problem




3 different moving stripes.

Observed from the hole, is same.

Thus direction of optical flow: **multiple direction**


unable to estimate direction of optical flow

Aperture problem




Aperture problem

Multi-direction



Aperture problem

single-direction



Conditions for solvability

- Optimal (u, v) satisfies Lucas-Kanade equation

$$\begin{matrix} \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} & \begin{bmatrix} u \\ v \end{bmatrix} & = & - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix} \\ A^T A & & & A^T b \end{matrix}$$

When is This **Solvable**? → No aperture problem


- $A^T A$ should be invertible
- $A^T A$ should not be too small due to noise
 - eigenvalues λ_1 and λ_2 of $A^T A$ should not be too small
- $A^T A$ should be well-conditioned
 - λ_1 / λ_2 should not be too large (λ_1 = larger eigenvalue)

$A^T A$ is solvable when there is no aperture problem


$$A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

Local Patch Analysis


Edge: multi-direction



Corner: single-direction



Aperture problem cause
multi-direction when at
edge, flat points



Flat: multi-direction



So need Corner detector!
Corner point: $A^T A$ Solvable


Eigenvectors of $A^T A$

$$A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$



- Recall the Harris corner detector: $M = A^T A$ is the *second moment matrix*
- The eigenvectors and eigenvalues of M relate to edge direction and magnitude
 - The eigenvector associated with the larger eigenvalue points in the direction of fastest intensity change
 - The other eigenvector is orthogonal to it

Interpreting the eigenvalues

Classification of image points using eigenvalues of the second moment matrix:



Edge






$$\sum \nabla I (\nabla I)^T$$

- large gradients, all the same
- large λ_1 , small λ_2



Low texture region





$$\sum \nabla I (\nabla I)^T$$

- gradients have small magnitude
- small λ_1 , small λ_2



High textured region/Corner



$$\sum \nabla I (\nabla I)^T$$

- gradients are different, large magnitudes
- large λ_1 , large λ_2



Horn-Schunck

Optical flow equation

Combining these two equations

$$\begin{aligned}0 &= I(x + u, y + v) - H(x, y) & I_x &= \frac{\partial I}{\partial x} \\&\approx I(x, y) + I_x u + I_y v - H(x, y) \\&\approx (I(x, y) - H(x, y)) + I_x u + I_y v \\&\approx I_t + I_x u + I_y v \\&\approx I_t + \nabla I \cdot [u \ v]\end{aligned}$$

In the limit as u and v go to zero, this becomes exact


$$0 = I_t + \nabla I \cdot \left[\frac{\partial x}{\partial t} \ \frac{\partial y}{\partial t} \right]$$

Horn-Schunck

Horn-Schunck algorithm: global method
global constraint of smoothness
solve the aperture problem

LK constraint: $\min_{u,v} \left[I_x u_{ij} + I_y v_{ij} + I_t \right]^2$

Smooth constraint:



Global energy function:

$$E = \iint \left[(I_x u + I_y v + I_t)^2 + \alpha^2 (\|\nabla u\|^2 + \|\nabla v\|^2) \right] dx dy$$

Horn-Schunck

Discrete global energy function:

$$E = \sum_{i,j} E_s(i,j) + \lambda \sum_{i,j} E_d(i,j)$$

$$E_s(i,j) = \frac{1}{4} [(u_{i,j} - u_{i+1,j})^2 + (u_{i,j} - u_{i,j+1})^2 + (v_{i,j} - v_{i+1,j})^2 + (v_{i,j} - v_{i,j+1})^2]$$

$$E_d(i,j) = \|I_x u_{ij} + I_y v_{ij} + I_t\|^2$$


I_x, I_y, I_t : derivatives of image intensity values along x, y, t

$\vec{V} = [u(x,y), v(x,y)]^\top$: optical flow vector (to be solved)

Horn-Schunck

Minimize discrete global energy function:


$$\sum_{ij} \left\{ \frac{1}{4} \left[(u_{ij} - u_{i+1,j})^2 + (u_{ij} - u_{i,j+1})^2 + (v_{ij} - v_{i+1,j})^2 + (v_{ij} - v_{i,j+1})^2 \right] + \lambda \left[I_x u_{ij} + I_y v_{ij} + I_t \right]^2 \right\}$$



$$(u_{ij}^2 - 2u_{ij}u_{i+1,j} + u_{i+1,j}^2)$$

$$(u_{ij}^2 - 2u_{ij}u_{i,j+1} + u_{i,j+1}^2)$$

(variable will appear four times in sum)



$$\frac{\partial E}{\partial u_{kl}} = 2(u_{kl} - \bar{u}_{kl}) + 2\lambda(I_x u_{kl} + I_y v_{kl} + I_t)I_x$$

$$\frac{\partial E}{\partial v_{kl}} = 2(v_{kl} - \bar{v}_{kl}) + 2\lambda(I_x u_{kl} + I_y v_{kl} + I_t)I_y$$

short hand for
local average

$$\bar{u}_{ij} = \frac{1}{4} \left\{ u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1} \right\}$$

Horn-Schunck

Minimize discrete global energy function:

$$\frac{\partial E}{\partial u_{kl}} = 2(u_{kl} - \bar{u}_{kl}) + 2\lambda(I_x u_{kl} + I_y v_{kl} + I_t)I_x$$

$$\frac{\partial E}{\partial v_{kl}} = 2(v_{kl} - \bar{v}_{kl}) + 2\lambda(I_x u_{kl} + I_y v_{kl} + I_t)I_y$$

(set derivatives to zero and solve for unknowns u and v)

$$(1 + \lambda I_x^2)u_{kl} + \lambda I_x I_y v_{kl} = \bar{u}_{kl} - \lambda I_x I_t$$

$$\lambda I_x I_y u_{kl} + (1 + \lambda I_y^2)v_{kl} = \bar{v}_{kl} - \lambda I_y I_t$$

this is a linear system

$$\mathbf{Ax} = \mathbf{b}$$

how do you solve this?

<https://blog.csdn.net/platane>

Horn-Schunck

Minimize discrete global energy function:

Matrix form:

$$\begin{bmatrix} 1 + \lambda I_x^2 & \lambda I_x I_y \\ \lambda I_x I_y & 1 + \lambda I_y^2 \end{bmatrix} \begin{bmatrix} u_{kl} \\ v_{kl} \end{bmatrix} = \begin{bmatrix} \bar{u}_{kl} - \lambda I_x I_t \\ \bar{v}_{kl} - \lambda I_y I_t \end{bmatrix}$$



$$\{1 + \lambda (I_x^2 + I_y^2)\} u_{kl} = (1 + \lambda I_y^2) \bar{u}_{kl} - \lambda I_x I_y \bar{v}_{kl} - \lambda I_x I_t$$

$$\{1 + \lambda (I_x^2 + I_y^2)\} v_{kl} = (1 + \lambda I_x^2) \bar{v}_{kl} - \lambda I_x I_y \bar{u}_{kl} - \lambda I_y I_t$$



$$u_{kl} = \bar{u}_{kl} - \frac{I_x \bar{u}_{kl} + I_y \bar{v}_{kl} + I_t}{\lambda^{-1} + I_x^2 + I_y^2} I_x, \quad v_{kl} = \bar{v}_{kl} - \frac{I_x \bar{u}_{kl} + I_y \bar{v}_{kl} + I_t}{\lambda^{-1} + I_x^2 + I_y^2} I_y$$

Horn-Schunck

Algorithm flow:

- Initialize the optical flow, $u = 0, v = 0$
- Iteration:
 - Calculate the image gradient I_x, I_y
 - Calculate the difference between the two frames I_t
 - Solve the below two equations iteratively until convergence.

$$u_{kl} = \bar{u}_{kl} - \frac{I_x \bar{u}_{kl} + I_y \bar{v}_{kl} + I_t}{\lambda^{-1} + I_x^2 + I_y^2} I_x, \quad v_{kl} = \bar{v}_{kl} - \frac{I_x \bar{u}_{kl} + I_y \bar{v}_{kl} + I_t}{\lambda^{-1} + I_x^2 + I_y^2} I_y$$

$$\bar{u}_{ij} = \frac{1}{4} \left\{ u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1} \right\}$$

Optical flow pyramid


Optical flow pyramid

Both the LK and Horn-schunck algorithms have the assumption of small motion.

However, this assumption cannot be guaranteed under actual scenarios.

Coarse-to-fine optical flow estimation

Optical flow pyramid algorithm to improve the drawbacks associated with the small motion assumption.




Gaussian pyramid of image H

$u=1.25$ pixels

$u=2.5$ pixels


$u=5$ pixels

$u=10$ pixels



Gaussian pyramid of image I

Coarse-to-fine optical flow estimation



Optical flow pyramid

Generate layer image:

- Image1, image2 are scaled by a certain ratio
- Bottom layer is the original size image
- The algorithm calculates the optical flow from the top layer
- Result of the previous layer as next layer input

Therefore this process is also called coarse-to-fine.

Optical flow pyramid

Generate layer image:

$$\begin{aligned} I^L(x, y) = & \frac{1}{4} I^{L-1}(2x, 2y) + \\ & \frac{1}{8} (I^{L-1}(2x-1, 2y) + I^{L-1}(2x+1, 2y) + I^{L-1}(2x, 2y-1) + I^{L-1}(2x, 2y+1)) + \\ & \frac{1}{16} (I^{L-1}(2x-1, 2y-1) + I^{L-1}(2x+1, 2y+1) + I^{L-1}(2x-1, 2y+1) + I^{L-1}(2x+1, 2y-1)). \end{aligned}$$

I_0 : Layer 0 of image I , the original image with the highest resolution

I^{L-1} : L define number of layers. I^{L-1} denotes the image of layer $L-1$


$$\mathbf{g}^{L-1} = 2(\mathbf{g}^L + \mathbf{d}^L).$$

I^{L-1} : g^{L-1} denotes the initial optical flow of layer $L-1$.

d^L denotes the estimated optical flow of the L layer.

Optical flow pyramid

Generate layer image:



Initial optical flow layer, as the initial input for the next layer.

Using LK or HS iterative optical flow for each layer.

Techniques for estimate motion

Feature-based methods (e.g. SIFT+Ransac+regression)

- Extract visual features (corners, textured areas) and track them over multiple frames
- Sparse motion fields, but possibly robust tracking
- Suitable especially when image motion is large (10-s of pixels)

Direct-methods (e.g. optical flow)

- Directly recover image motion from spatio-temporal image brightness variations
- Global motion parameters directly recovered without an intermediate feature motion calculation
- Dense motion fields, but more sensitive to appearance variations
- Suitable for video and when image motion is small (< 10 pixels)